



# APRENDIZAJE REFORZADO

## CLASE 6

**Julían Martínez**

# SARSA: TD CON ON-POLICY CONTROL



Sarsa

$$\hat{Q}^{t+1}(S_t, A_t) = \hat{Q}^t(S_t, A_t) + \alpha [R_{t+1} + \gamma \hat{Q}^t(S_{t+1}, A_{t+1}) - \hat{Q}^t(S_t, A_t)]$$

$$A_{t+1} \sim \pi_{\hat{Q}^t} - \epsilon \text{ greedy}$$

# PSEUDO-CÓDIGO DEL SUTTON

## Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R$ ,  $S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$



$S \leftarrow S'$ ;  $A \leftarrow A'$ ;

    until  $S$  is terminal

# WINDY GRIDWORLD




El reward es -1  
hasta llegar a G



Intensidad del viento

$\epsilon$  - GREEDY SARSA,  $\epsilon = 0.1$ ,  $\alpha = 0.5$



## Q-LEARNING: TD CON OFF-POLICY CONTROL

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right].$$

## Q-LEARNING: TD CON OFF-POLICY CONTROL

$$\hat{q}^{t+1}(S_t, A_t) = \hat{q}^t(S_t, A_t) + \alpha [R_{t+1} + \gamma \hat{q}^t(S_{t+1}, A_{t+1}^*) - \hat{q}^t(S_t, A_t)]$$

$$A_{t+1}^* = \pi_{\hat{q}^t}^*(S_{t+1}) \quad , \quad A_{t+1} \sim \pi_{\hat{q}^t}^{*, \epsilon}(\cdot | S_{t+1})$$

# PSEUDO-CÓDIGO DEL SUTTON

## Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$


$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$


    until  $S$  is terminal



## EJEMPLO 6.6 - SUTTON (ACANTILADO WALKING)



# SARSA VS Q-LEARNING



¿POR QUÉ ES OFF-POLICY?

$$q_*(s, a) = \mathbb{E}_{\pi_*} [G_t \mid S_t = s, A_t = a]$$

$$= \mathbb{E}_{\pi_*} [R_{t+1} + \gamma q_{\pi_*}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$

$$= \mathbb{E}_{\pi_*^\epsilon} [\{R_{t+1} + \gamma q_{\pi_*}(S_{t+1}, A_{t+1})\} w_{t+1:T} \mid S_t = s, A_t = a]$$

$$w(s', a') = \frac{1_{\{\pi^*(s') = a'\}}}{(1-\epsilon) 1_{\{\pi^*(s') = a'\}} + \epsilon \cdot \frac{1}{|A|}}$$

$\mathcal{L} = \phi \psi^{\dagger}$   
 $\nabla \cdot \mathbf{E} = 0 \quad \frac{\partial \mathbf{H}}{\partial t} = -\nabla \phi$   
 $\nabla \cdot \mathbf{H} = 0 \quad \frac{\partial \mathbf{E}}{\partial t} = -\nabla \psi$   
 $\frac{\partial \psi}{\partial t} = \psi$   
 $f(w) = \int_{-\infty}^{\infty} f(x) e^{-i w \cdot x} dx \quad \frac{d}{d\theta}$   
 $\rho \left( \frac{\partial^2}{\partial t^2} + \mathbf{v} \cdot \nabla \right) = -\nabla p + \nabla T + \mathbf{f}$   
 $H = - \sum_{i=1}^n p(x) \log p(x)$   
 $\frac{1}{2} G^2 S^2 \frac{\partial^2 V}{\partial S^2} + r S \frac{\partial V}{\partial S} + \frac{\partial V}{\partial t} - r \cdot V = 0$   
 $(Q, q, m) = \sum_{i=1}^n \left[ \frac{D_i}{m_i q_i} S_{i-1} + c_i D_i + \frac{q_i H_i}{2} \left( m_i \left( 1 - \frac{D_i}{p_i} \right) - 1 + \frac{1}{2} \right) \right]$   
 $\left[ \begin{array}{c} \frac{d \Delta p(s, \phi)}{d \phi} \\ \frac{d \Delta M(s, \phi)}{d \phi} \end{array} \right] = \left[ \begin{array}{cc} \beta & -\beta \\ -\beta & 0 \end{array} \right] \left[ \begin{array}{c} \Delta p(s, \phi) \\ \Delta M(s, \phi) \end{array} \right]$   
 $\int_0^{\pi} (\log \sin x)^2 dx = - \int_0^{\pi} (\log \cos x)^2 dx = - \frac{\pi}{2} \left\{ \frac{\pi^2}{12} + (\log 2)^2 \right\}$

## RESUMIENDO

$$\hat{q}^{t+1}(S_t, A_t) = \hat{q}^t(S_t, A_t) + \alpha [R_{t+1} + \gamma \hat{q}^t(S_{t+1}, A_{t+1}^*) - \hat{q}^t(S_t, A_t)]$$

$$A_{t+1}^* = \pi_{\hat{q}^t}^*(S_{t+1}) \quad , \quad A_{t+1} \sim \pi_{\hat{q}^t}^{*, \epsilon}(\cdot | S_{t+1})$$

$$= \mathbb{E}_{\pi_*^\epsilon} \left[ \left\{ R_{t+1} + \gamma q_{\pi_*}(S_{t+1}, A_{t+1}) \right\} w(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a \right]$$

$$w(s', a') = \frac{1_{\{\pi^*(s') = a'\}}}{(1-\epsilon) 1_{\{\pi^*(s') = a'\}} + \epsilon \cdot \frac{1}{|A|}}$$

# STOCHASTIC APPROXIMATION

$$E[f(\theta, W)]|_{\theta=\theta^*} = 0$$

Robbins-Monro, 1951

$$\theta(n+1) = \theta(n) + \alpha_n f(\theta(n), W(n+1))$$

$$\sum \alpha_n = \infty$$

$$\sum \alpha_n^2 < \infty$$

$$\alpha_n = \frac{1}{n}$$

# ODES / MÉTODO DE EULER

$$\bar{f}(\theta) = E[f(\theta, w)]$$

$$\frac{dx(t)}{dt} = \bar{f}(x(t))$$

$$\frac{\theta(n+1) - \theta(n)}{\frac{1}{n}} = \bar{f}(\theta(n))$$

$$\theta(n+1) = \theta(n) + \alpha_n [\bar{f}(\theta(n)) + \Delta(n+1)]$$

$$\Delta(n+1) = f(\theta(n), W(n+1)) - \bar{f}(\theta(n))$$

THM If the ODE has a unique asymptotically stable equilibrium  $x^* \Rightarrow x_n \rightarrow x^*$  with probability one

$$P\left(\lim_{n \rightarrow \infty} x_n = x^*\right) = 1$$



EJEMPLO 1 - MONTE CARLO

$$E[C(X)] \stackrel{?}{=}$$

$$f(\theta, x) = C(x) - \theta$$

$$\theta_{n+1} = \theta_n + \alpha_n [C(X_{n+1}) - \theta_n]$$

EJEMPLO 2 - AJUSTE

$$F(\theta) := \frac{1}{2} E[(Y - f_{\theta}(X))^2]$$

$$\nabla_{\theta} F(\theta) = E[(Y - f_{\theta}(X)) \cdot \nabla_{\theta} f_{\theta}(X)]$$

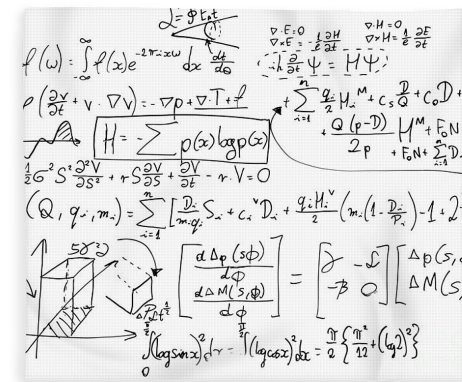
$$\theta_{n+1} = \theta_n + \alpha_n [(Y_n - f_{\theta_n}(X_n)) \nabla_{\theta} f_{\theta_n}(X_n)]$$

### EJEMPLO 3 - TEMPORAL DIFFERENCE

$$E\left[R_s^+ + \gamma \sum_{s'} v(s') \mathbb{1}_{\{S_s^+ = s'\}} - v(s)\right] = 0 \quad \forall s$$


$$v_{n+1} = v_n + \alpha \left[ R_{n+1} + \gamma v_n(S_{n+1}) - v_n(S_n) \right]$$

# ELIGIBILITY TRACES



$$E_0(s) = 0$$

$$E_t(s) = \gamma \lambda E_{t-1}(s) + \mathbf{1}(S_t = s)$$



accumulating eligibility trace

times of visits to a state

# EL LADO OSCURO...



- Artículo de divulgación sobre Stochastic Approximation  
<https://www.ias.ac.in/article/fulltext/reso/018/12/1086-1094>
- THE O.D.E. METHOD FOR CONVERGENCE OF STOCHASTIC APPROXIMATION AND REINFORCEMENT LEARNING  
<http://repository.ias.ac.in/5333/1/351.pdf>
- Reinforcement Learning: Hidden Theory and New Super-Fast Algorithms (Charla de Meyn sobre Stochastic Approximation en el Simons Institute) <https://www.youtube.com/watch?v=dhEF5pfYmvc>

